

Floras Sales Intelligence Demo

Krishna Narayanasamy

Fill Example: Patagonia

Company name

Industry

Website
https://www.patagonia.com"/>

Notes


Add anything you know about the company, sustainability goals, emissions, or customer base.

Generate Sales Intelligence

Demo Preview

Overview

- **What it is:** A proof-of-concept Floras sales intelligence tool. It takes a company name, industry, website, and notes, then generates a tailored report with entity type, financial profile, fit score, matched use case, outreach pitch, risks, and quality review.
- **How it works:** The system combines code-based lead research with GPT agents. It scrapes company websites, searches ESG/business signals with Tavily, checks SEC 10-K and Wikipedia when available, and uses notes as extra context. Rules classify the company as buyer, supplier, bank, or other; GPT refines ambiguous cases. Optional FMP data adds margins and ROA for public companies.
- **Floras intelligence layer:** The tool uses Floras overview knowledge, entity-type rules, and 12 structured sales cases covering corporate travel/SAF, CPG procurement, supplier differentiation, logistics, CSRD disclosure, credit cards, financial services, and more. The best-matching case shapes the use case, pitch, safe claims, and project suggestions.
- **Output:** A report with company summary, hierarchy, financial fit, High/Medium/Low score, key sales signals, personalized Floras use case, suggested project types, outreach pitch, research sources, matched case badge, risks/unknowns, and reviewer suggestions.
- **Stack:** Plain HTML frontend, FastAPI backend, OpenAI, Tavily, optional FMP.



Sales Intelligence Steps

1

Enter the
Target

Company name,
industry, website,
notes

2

Lead
Research
Agent

Website + SEC +
Wikipedia + Tavily
ESG; classify buyer
/ supplier / bank,
research cache

3

Floras
knowledge
matching

KB + 12 sales
cases + financial
profile (FMP) →
best-fit playbook

4

Proposal
Agent

GPT drafts fit score,
use case, pitch
from research +
matched case

5

Quality review
& report

GPT checks
grounding, entity fit,
financials; outputs
final BD report

Lead Research Agent

- **First pipeline step** (lead_research.py) — gathers raw company intelligence **before** any report is drafted.
- **Mostly code, not AI** — HTTP, HTML parsing, Jina Reader, Tavily, SEC EDGAR, Wikipedia; optional GPT only for entity refinement later.
- **Entity preview** — rules classify buyer / supplier / financial issuer / poor fit / dual role from name + industry + notes → tunes scrape paths and Tavily queries.
- **Website scrape** (if URL): entity-specific paths; Jina first, HTML fallback; skips bot-gated or thin pages.
- **Tavily ESG search** — sustainability/climate context when company name is provided; query focus varies by entity type.
- **Business profile** (business_research.py): SEC 10-K Item 1 Business (US public) → Wikipedia → conditional Tavily business search (only if free sources are thin).
- **User notes** appended as a trusted source; output = research context, research_urls, source label, preview, and **entity_classification**.

Business Entity Classification

- Answers **where the lead sits in the value chain** — wrong entity type means the wrong Floras pitch (wallet vs supplier embed vs financial product).
- **Five types:** Enterprise buyer, Supplier to enterprises, Financial product issuer, Poor fit entity, Dual buyer & supplier.
- **Rules-first** (entity_types.json + entity_type.py) — industry signals, text signals, and negative signals scored on name, industry, notes, and full research text.
- **Entity preview** runs before scraping to route paths and Tavily; **final classification** runs after research with optional **GPT refinement** when confidence is low or types tie.
- Each type defines **pitch focus**, **outreach target** (e.g. CPO vs supplier sales), **preferred/discouraged case IDs**, and research config.
- Feeds **case matching** (+3 boost / -4 penalty), **Proposal Agent** synthesis, and **Quality Reviewer** entity-alignment checks.
- **UI:** Business Hierarchy section — type, confidence, hierarchy level, pitch focus, outreach target, matched signals, LLM rationale if used.

Floras KB & Matching

- **Floras KB** is server-side reference material so outputs are Floras-specific, not generic sustainability chat.
- **Layer 1 — `floras_overview.txt`:** climate-fintech, transaction-linked climate value, enterprise wallet, supplier embed, ICP, fit guidance — included in **every** analysis.
- **Layer 2 — `entity_types.json`:** five entity types with scrape paths, Tavily focus, and case preferences — drives classification and matching.
- **Layer 3 — `kb/cases/*.json`: 12 GTM playbooks** (CPG procurement, corporate travel/SAF, supplier differentiation, logistics, CSRD, credit card, SaaS, financial low-fit, green loan, fashion, beauty, organic waste).
- **Case matching** (`kb.py`): scores triggers $\times 2$ + pain signals + sectors on research + industry + notes; min score 2; entity type boosts preferred cases (+3) and penalizes discouraged ones (−4).
- **Financial profile** (optional `FMP_API_KEY`): gross margin, operating margin, ROA for public companies — shown as Financial Profile in UI; fed into proposal and review prompts.
- **Not a live catalog** — static demo files; no connection to Floras verified project database.

Proposal Agent

- **First AI step** in main.py — OpenAI **gpt-4o-mini** turns research + Floras context into a structured sales intelligence report.
- **Inputs:** lead research, entity classification, financial profile (if available), Floras overview, matched sales case (if any), and **proposal_instructions.py** synthesis block.
- **Synthesis order:** (1) entity type → WHO to pitch & which Floras mechanism → (2) research facts → (3) financial plausibility → (4) matched case pitch structure → (5) output field rules.
- Prompted to stay **grounded** in research, use Floras terminology (wallet, supplier embed, certificates — not generic “carbon offset”), and flag thin research in risks_or_unknowns.
- **Fit score** (0–100) and level (High 75–100 / Medium 45–74 / Low 0–44) — AI-judged using entity type, research signals, financial margins, and Floras ICP.
- **Outreach pitch:** 1–2 sentences without a matched case; 2–4 short paragraphs with a matched case; `enrich_outreach_pitch()` may append case template content.
- **Returns JSON:** summary, fit, signals, use case, project types, pitch, risks — plus matched case metadata for the UI.

Quality Reviewer

- **Second AI step** (quality_reviewer.py) — OpenAI call #2, **gpt-4o-mini**, temperature **0.3** for a stricter audit.
- Runs **after** the proposal is generated and the outreach pitch is enriched — critiques the draft; does **not** rewrite the full report.
- **Inputs:** proposal JSON, research sample, Floras KB overview, **entity classification**, and **financial profile** (when available).
- **Checks:** company-specific vs generic use case; claims grounded in research vs invented; pitch aligned to entity type (buyer vs supplier vs bank).
- **Checks (continued):** Floras mechanisms used correctly; fit score justified vs ICP; gross margin reflected in fit/pitch when financials exist (implication, not stat dump).
- **Returns JSON:** quality_rating (Good / Fair / Needs Improvement), quality_review (2–4 sentences), review_improvements (actionable fixes).
- **UI:** shown at the bottom of the report so reps see trustworthiness and what to fix before customer-facing use.
- **Purpose:** separates **drafting** from **auditing** — catches generic language and Floras misalignment before the output is treated as final.

Significant Updates/Challenges

- **Added entity classification** — buyer / supplier / bank / poor fit / dual; routes research, cases, and pitch.
- **Added business-model research** — SEC 10-K, Wikipedia, conditional Tavily; not just sustainability scraping.
- **Added financial profile (FMP API)** — gross margin, operating margin, ROA for public companies (BD's top signal).
- **Expanded to 12 sales cases** with entity-aware matching and unified synthesis instructions for the Proposal Agent.
- **Challenges:** API cost/latency, Wikipedia/SEC headers, FMP endpoint + ticker false positives, fit score still AI-estimated not rule-based.
- **UI:** Business Hierarchy + Financial Profile sections so reps see *why* the tool scored and pitched the way it did.

FLORAS SALES INTELLIGENCE DEMO